

Lecture 9: Data Storage and IO Models

Announcements

- Submission Project Part 1 tonight
 - Instructions on Piazza!
- PS2 due on Friday at 11:59 pm
 - Questions? Easier than PS1.
- Badgers Rule!

看这次
activity
for PS2.



Today's Lecture

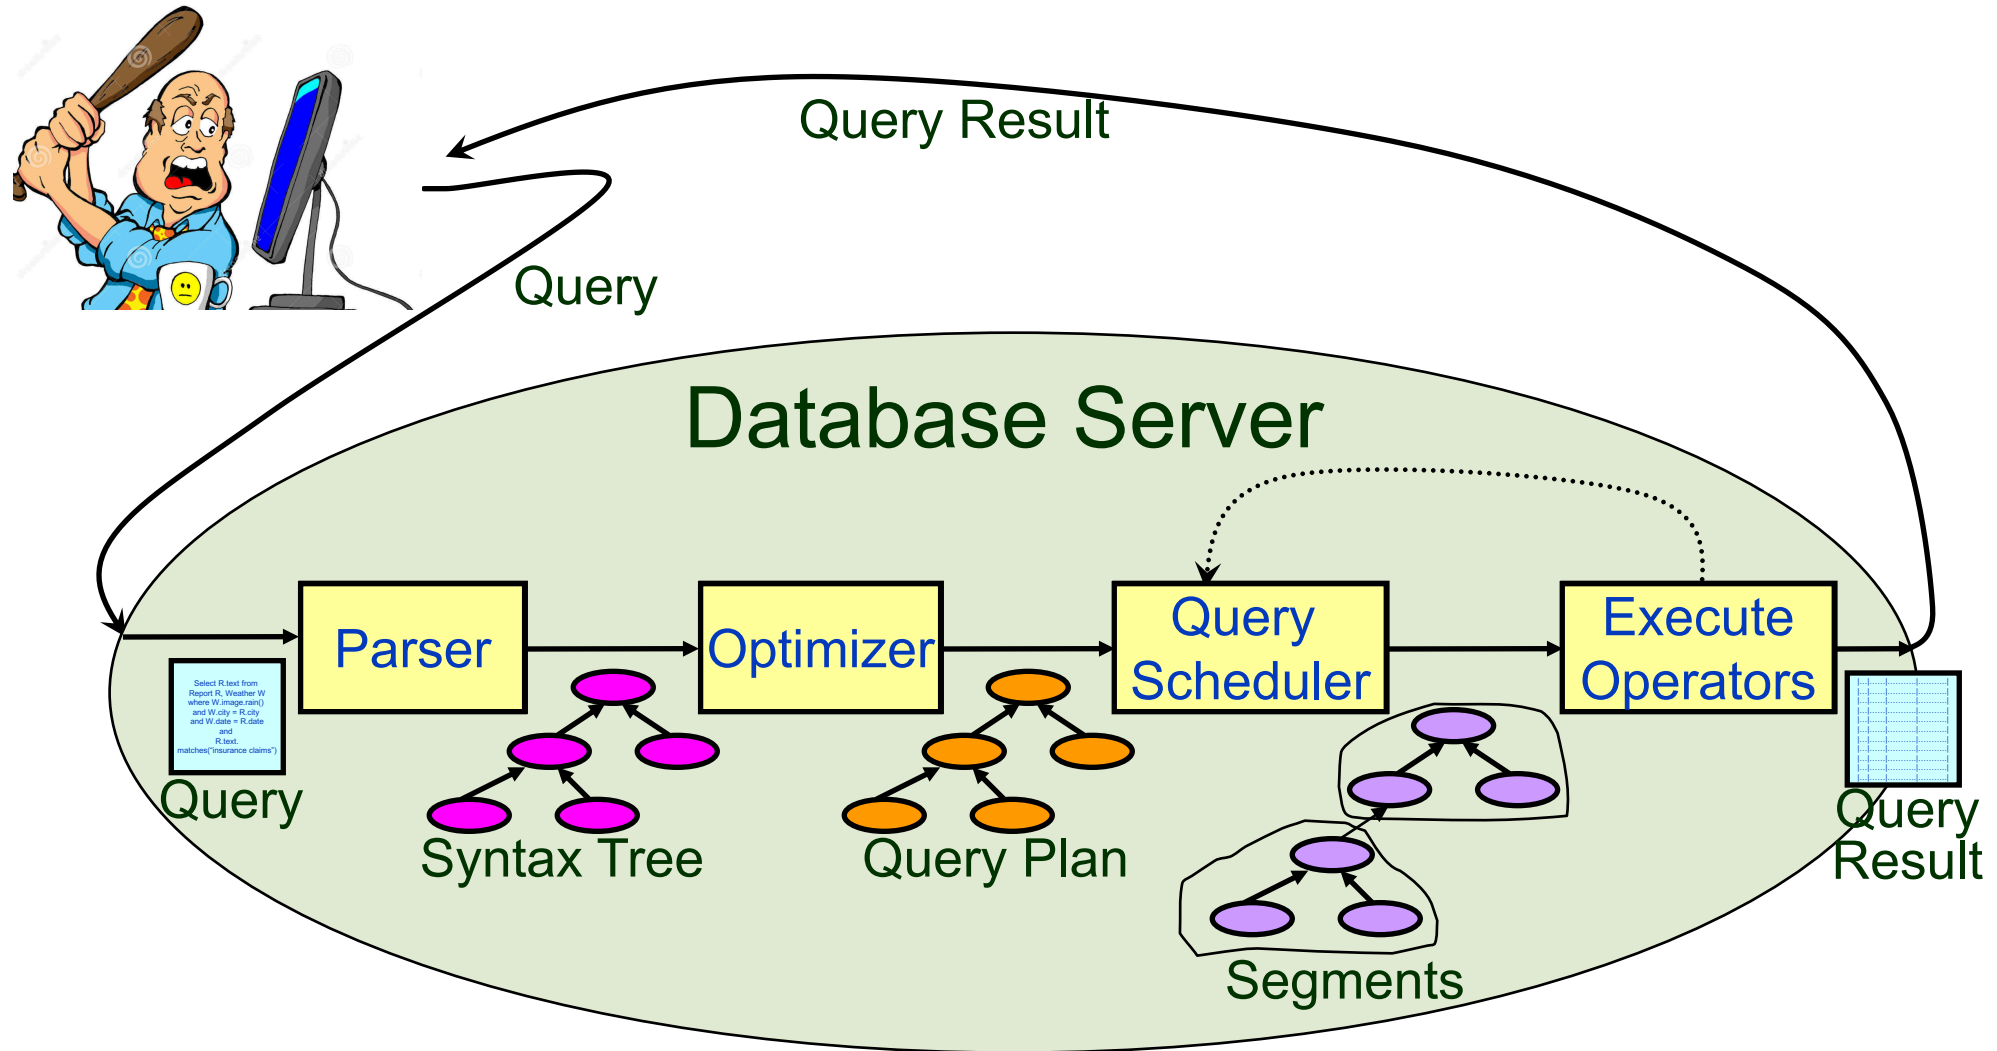
1. Data Storage
2. Disk and Files
3. Buffer Manager - Prelims

1. Data Storage

What you will learn about in this section

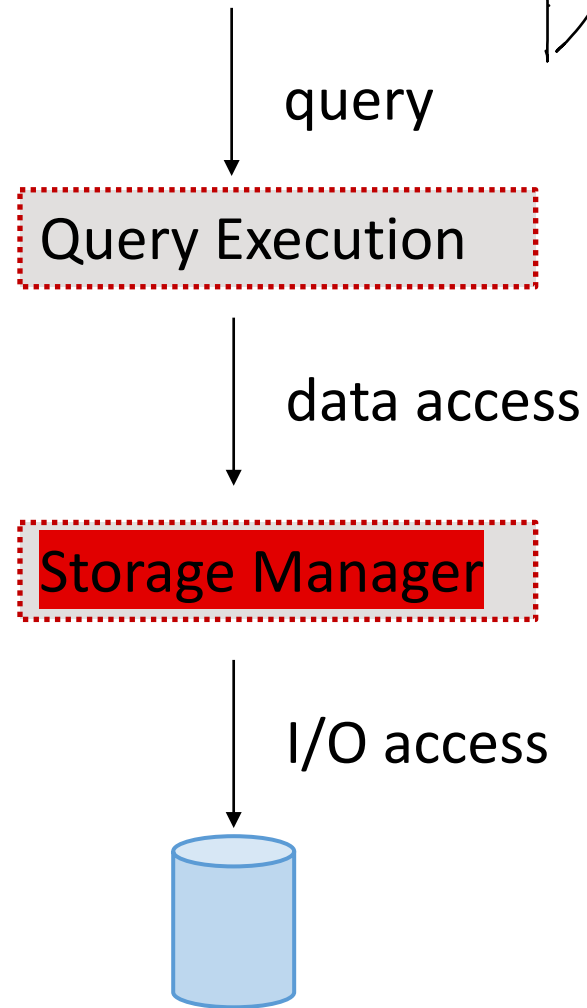
1. Life cycle of a query
2. Architecture of a DBMS
3. Memory Hierarchy

Life cycle of a query

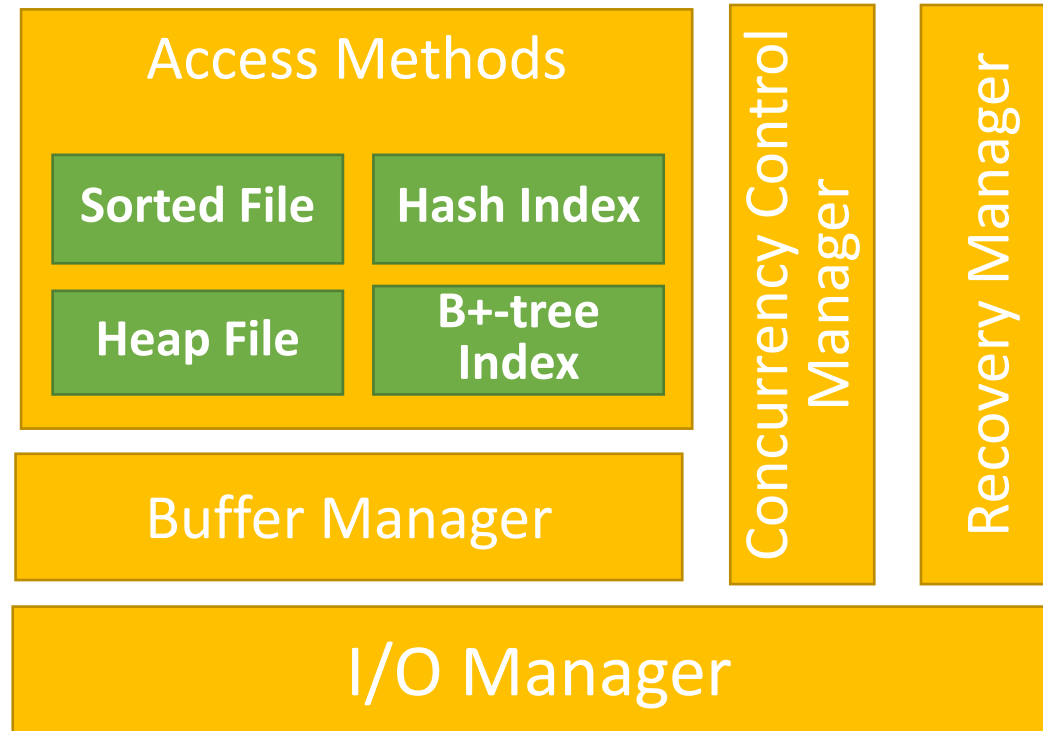


Internal Architecture of a DBMS

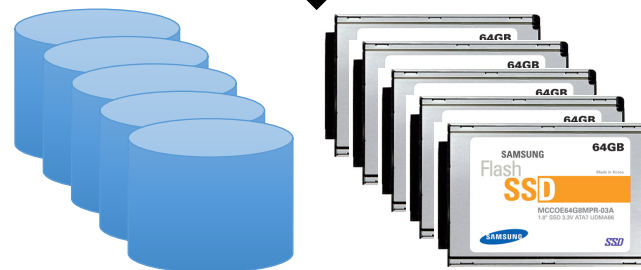
Database Management System.



Architecture of a Storage Manager



IO Accesses



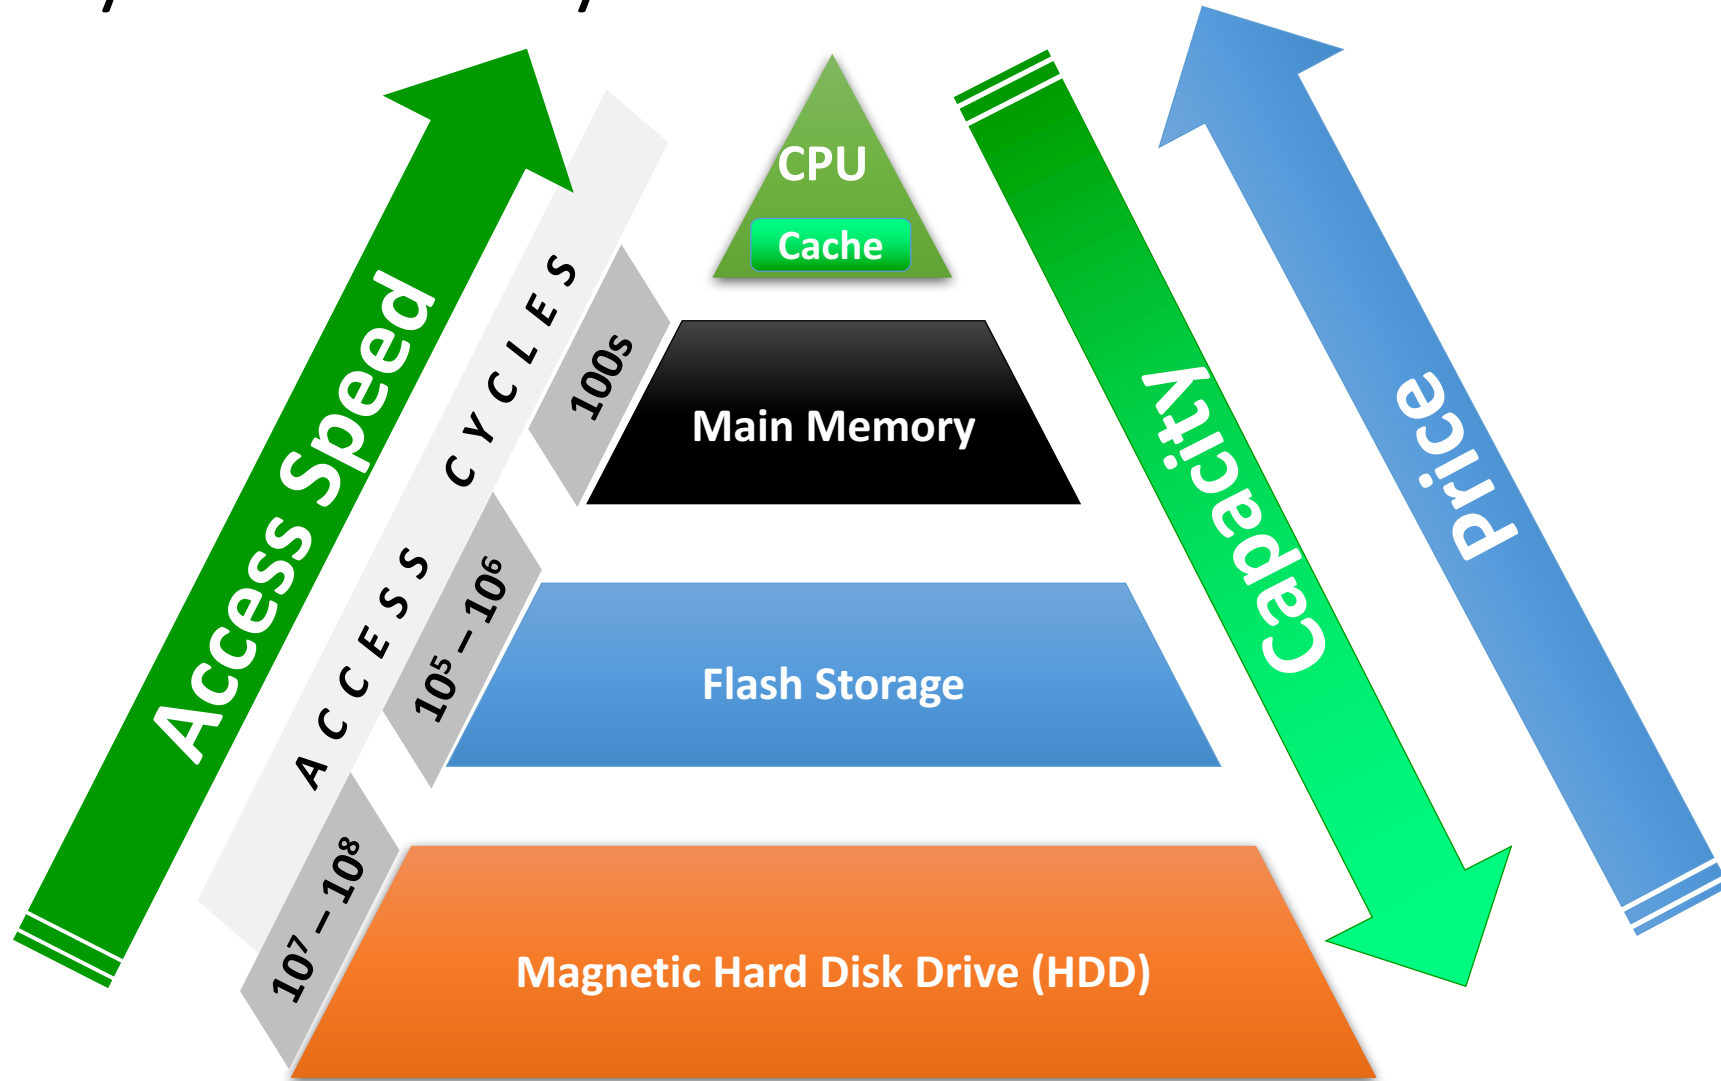
In Systems,
IO cost matters a ton!

Second part of project.

Data Storage

- How does a DBMS store and access data?
 - main memory (fast, temporary)
 - disk (slow, permanent)
- How do we move data from disk to main memory?
 - buffer manager
- How do we organize relational data into files?

Memory Hierarchy



Why not main memory?

- Relatively high cost
- Main memory is **not persistent!**
- Typical storage hierarchy:
 - **Primary storage:** **main memory** (RAM) for currently used data
 - **Secondary storage:** **disk** for the main database
 - **Tertiary storage:** **tapes** for archiving older versions of the data

2. Disk and Files

What you will learn about in this section

1. All about disks
2. Accessing a disk
3. Managing disk space

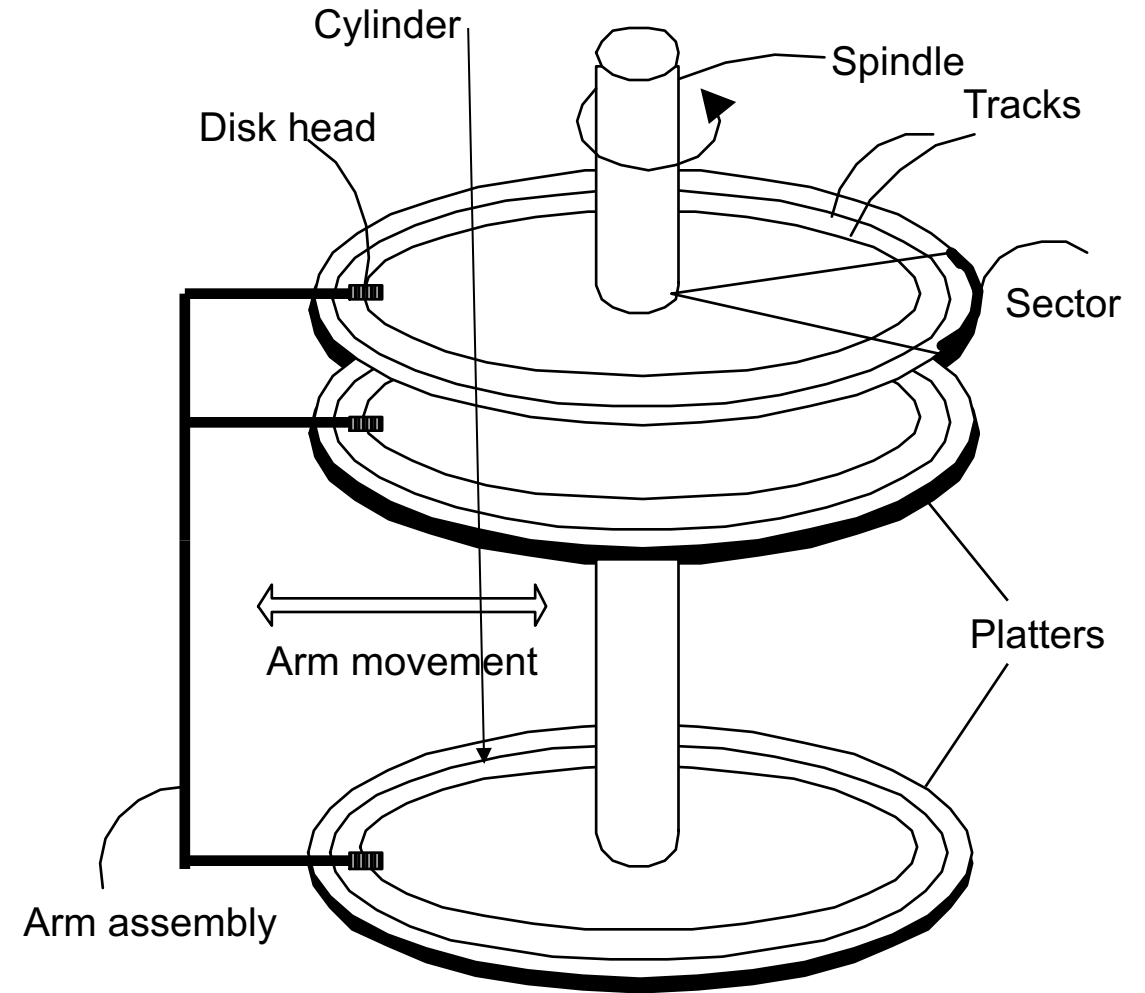
Disks

- Secondary storage device of choice.
- Data is stored and retrieved in units called *disk blocks*
- Unlike RAM, time to retrieve a disk page varies depending upon location on disk.
 - Therefore, relative placement of pages on disk has major impact on DBMS performance!

The Mechanics of a Disk

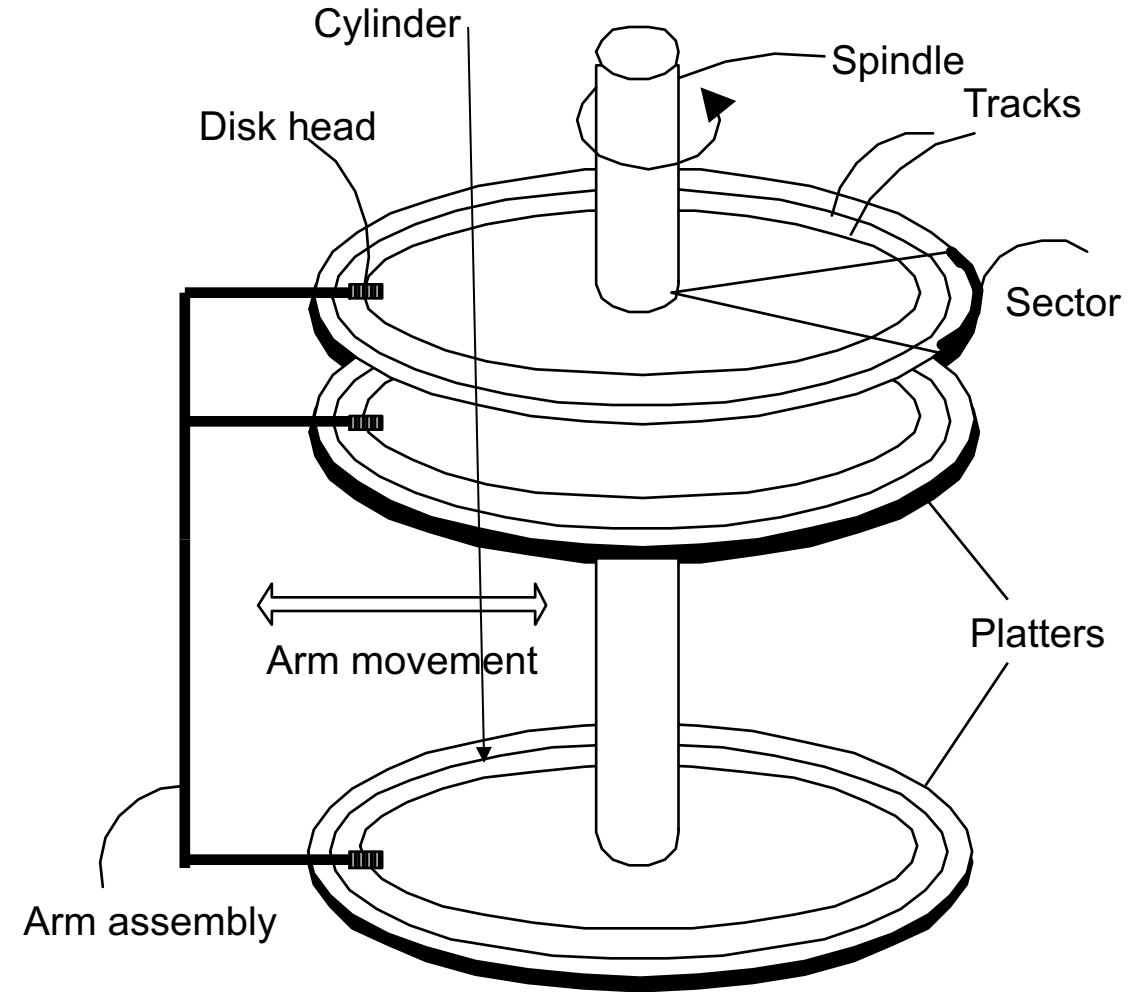
Mechanical characteristics:

- Rotation speed (7200RPM)
- Number of platters (1-30)
- Number of tracks (≤ 10000)
- Number of bytes/track (10^5)



The Mechanics of a Disk

- Platters spin @ ~ 7200rpm
- Arm assembly moves to position a head on a desired track. Tracks under heads make a **cylinder** (imaginary!)
- Only 1 head reads/writes at any time
- *Block size* : multiple of *sector size* (which is fixed).



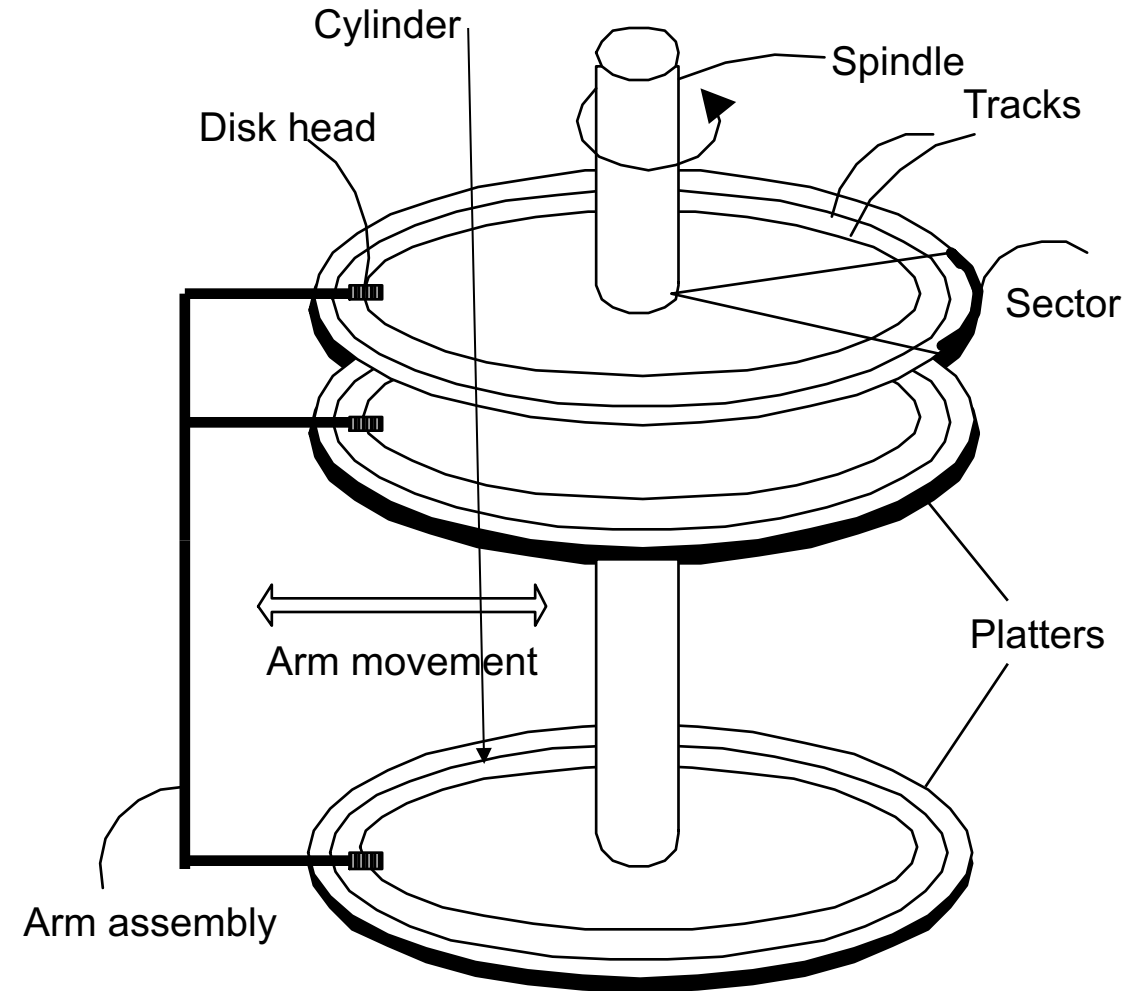
The Mechanics of a Disk

Unit of read or write:

disk block: $k \times \text{Sector Size}$

Once in memory: **page**

Typically: 4k or 8k or 16k



Access time = seek time + rotational delay + transfer time
(1-20 ms) (0-10ms) (~1 ms per 8k block)

The Mechanics of a Disk

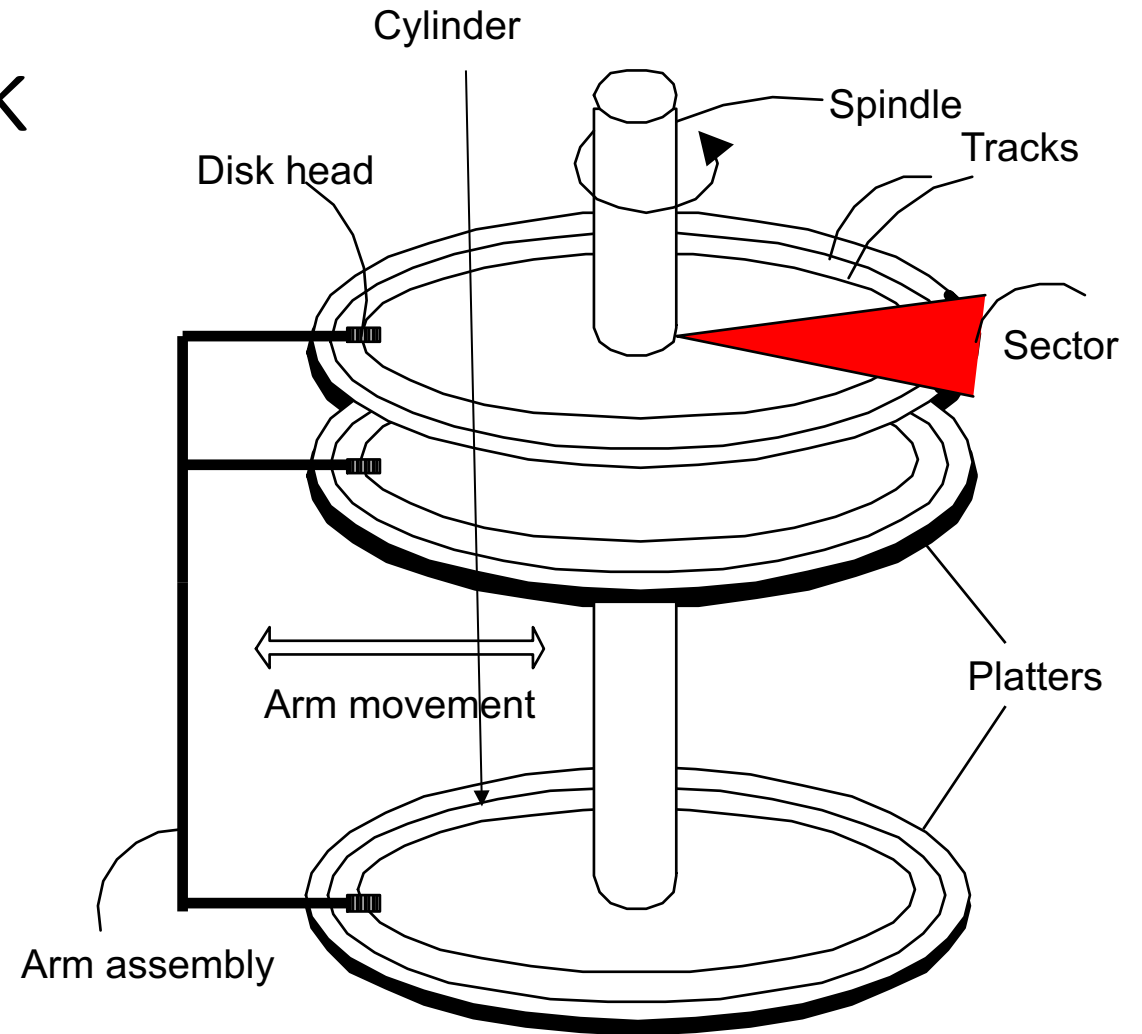
GOAL: Minimize seek and rotational delay

“Next Block” concept

- (1) Blocks on same track
- (2) Blocks on same cylinder
- (3) Blocks on adjacent cylinder

Disks read/write one block at a time

Access time = seek time +
rotational delay + transfer time



For a **sequential scan**, pre-fetching
several pages at a time is a big win!

Accessing the disk (I)

access time = **rotational delay** + seek time + transfer time

rotational delay: time to wait for sector to rotate under the disk head

- typical delay: *0–10 ms*
- maximum delay = 1 full rotation
- average delay ~ half rotation

RPM	Average delay
5,400	5.56
7,200	4.17
10,000	3.00
15,000	2.00

Accessing the disk (II)

access time = **rotational delay** + **seek time** + **transfer time**

seek time: time to move the arm to position disk head on the right track

- typical seek time: $\sim 9\text{ ms}$
- $\sim 4\text{ ms}$ for high-end disks

Accessing the disk (III)

access time = **rotational delay** + **seek time** + **transfer time**

data transfer time: time to move the data to/from the disk surface

- typical rates: $\sim 100 \text{ MB/s}$
- the access time is dominated by the seek time and rotational delay!

Example: Specs

	Seagate HDD
Capacity	3 TB
RPM	7,200
Average Seek Time	9 ms
Max Transfer Rate	210 MB/s
# Platters	3

What are the I/O rates for block size 4 KB and:

- random workload (~ 0.3 MB/s)
- sequential workload (~ 210 MB/s)

Managing Disk Space

- The disk space is organized into **files**
- Files are made up of **pages**
- Pages contain **records**

- Data is allocated/deallocated in increments of pages
- Logically close pages should be nearby in the disk

SSDs (Solid State Drive)

- SSDs use flash memory
- **No moving** parts (no rotate/seek motors)
 - eliminates seek time and rotational delay
 - very low power and lightweight
- Data transfer rates: 300-600 MB/s
- SSDs can read data (sequential **or** random) very fast!



SSDs

- Small storage (0.1-0.5x of HDD)
- expensive (20x of HDD)
- **Writes** are much more expensive than **reads** (10x)
- Limited lifetime
 - 1-10K writes per page
 - the average failure rate is 6 years

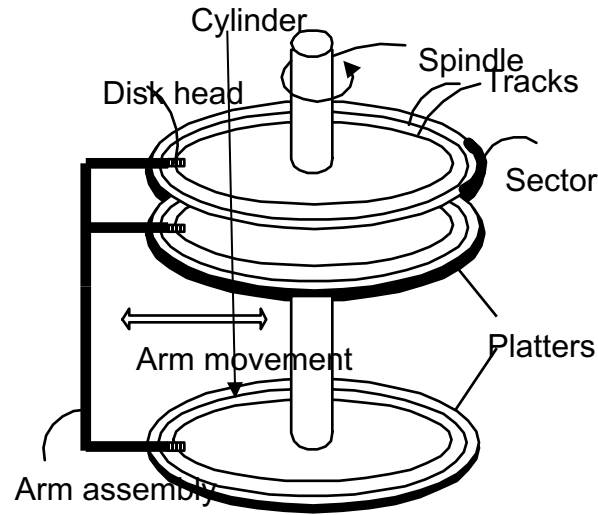
Can only read and write in blocks or pages of 2K, 4K, or more bytes. Looks like a disk.

3. Buffer Manager - Prelims

What you will learn about in this section

1. Buffer Manager
2. Replacement Policy

High-level: Disk vs. Main Memory



Disk:

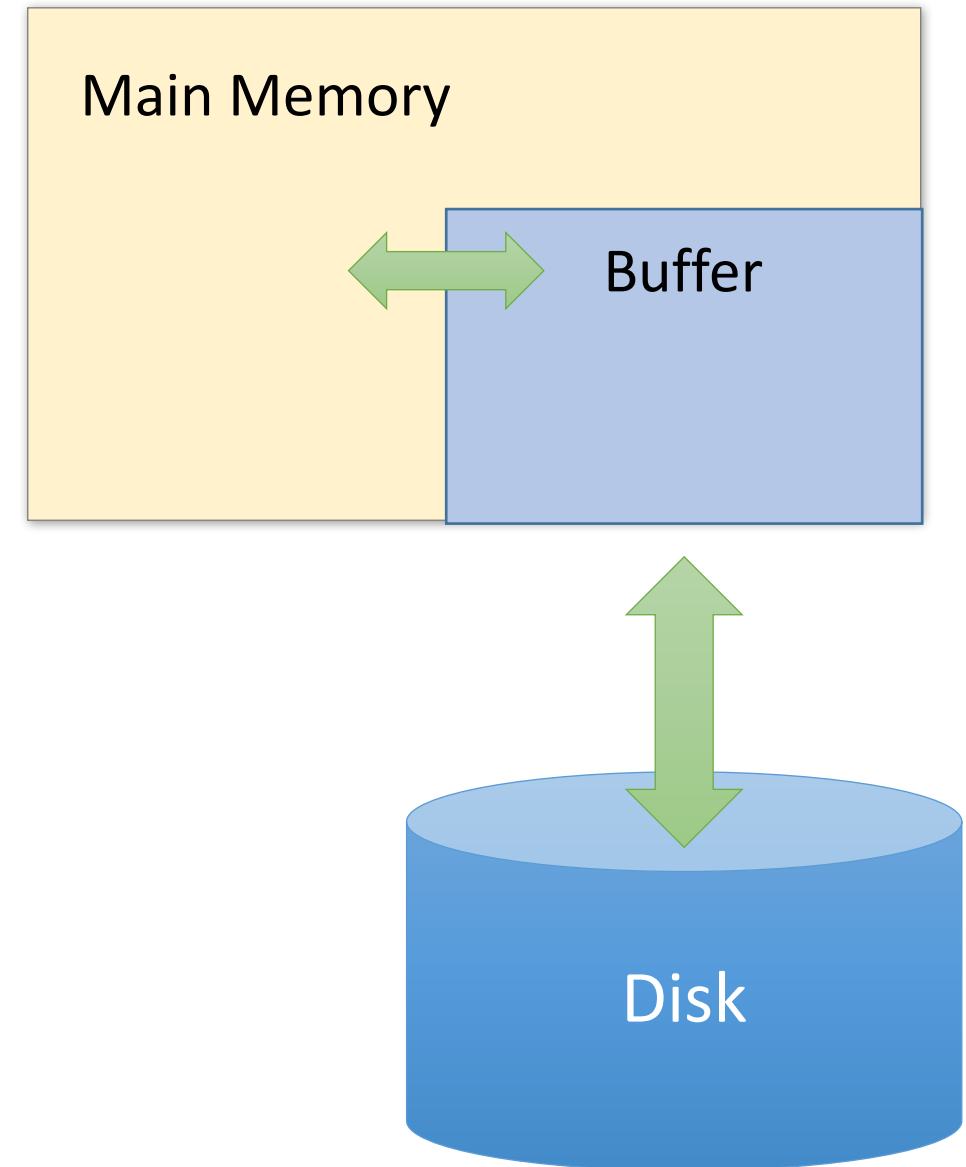
- **Slow:** Sequential *block* access
 - Read a blocks (not byte) at a time, so sequential access is cheaper than random
 - **Disk read / writes are expensive!**
- **Durable:** We will assume that once on disk, data is safe!
- **Cheap**

Random Access Memory (RAM) or Main Memory:

- **Fast:** Random access, byte addressable
 - ~10x faster for sequential access
 - ~100,000x faster for random access!
- **Volatile:** Data can be lost if e.g. crash occurs, power goes out, etc!
- **Expensive:** For \$100, get 16GB of RAM vs. 2TB of disk!

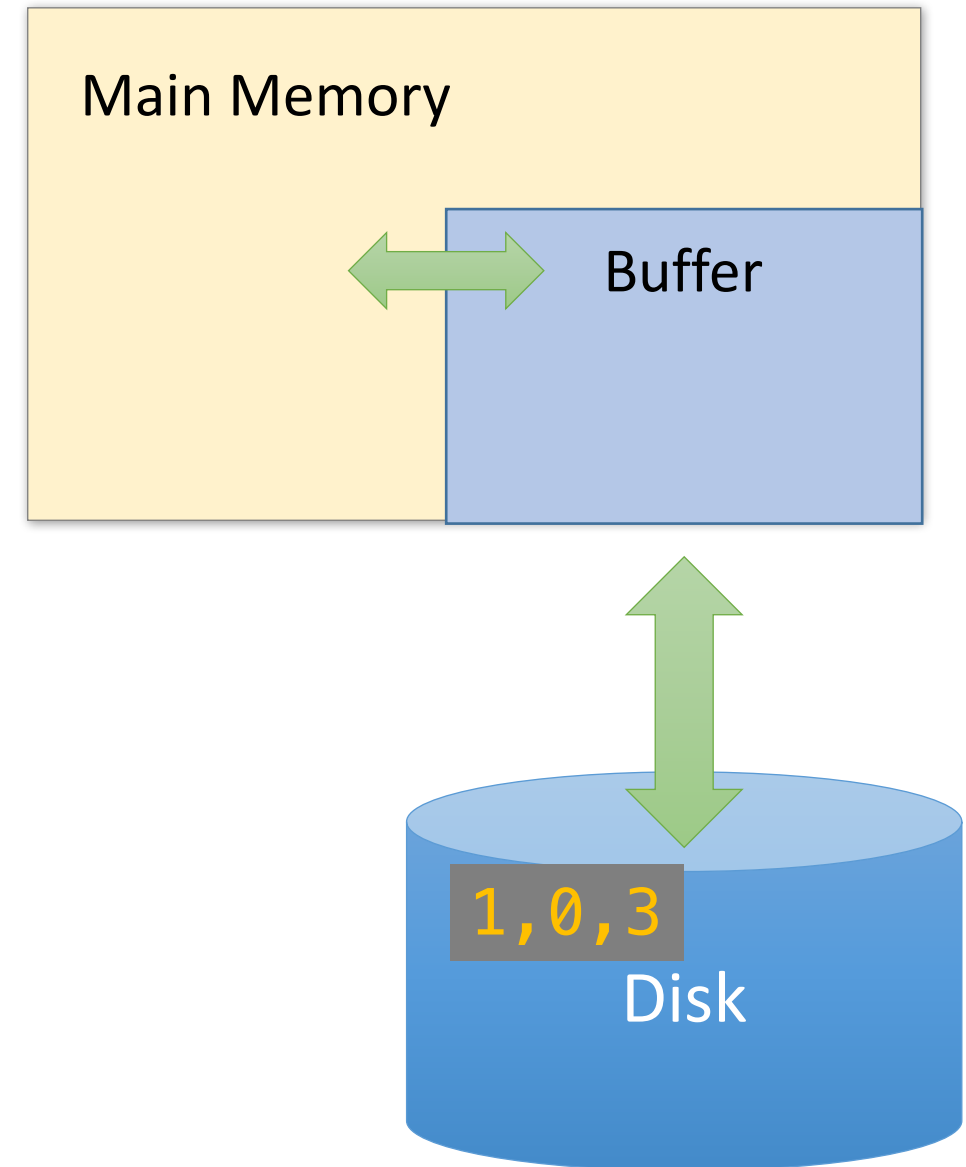
The Buffer

- A **buffer** is a region of physical memory used to store *temporary data*
 - *In this lecture:* a region in main memory used to store **intermediate data between disk and processes**
- *Key idea:* Reading / writing to disk is slow-need to cache data!



The (Simplified) Buffer

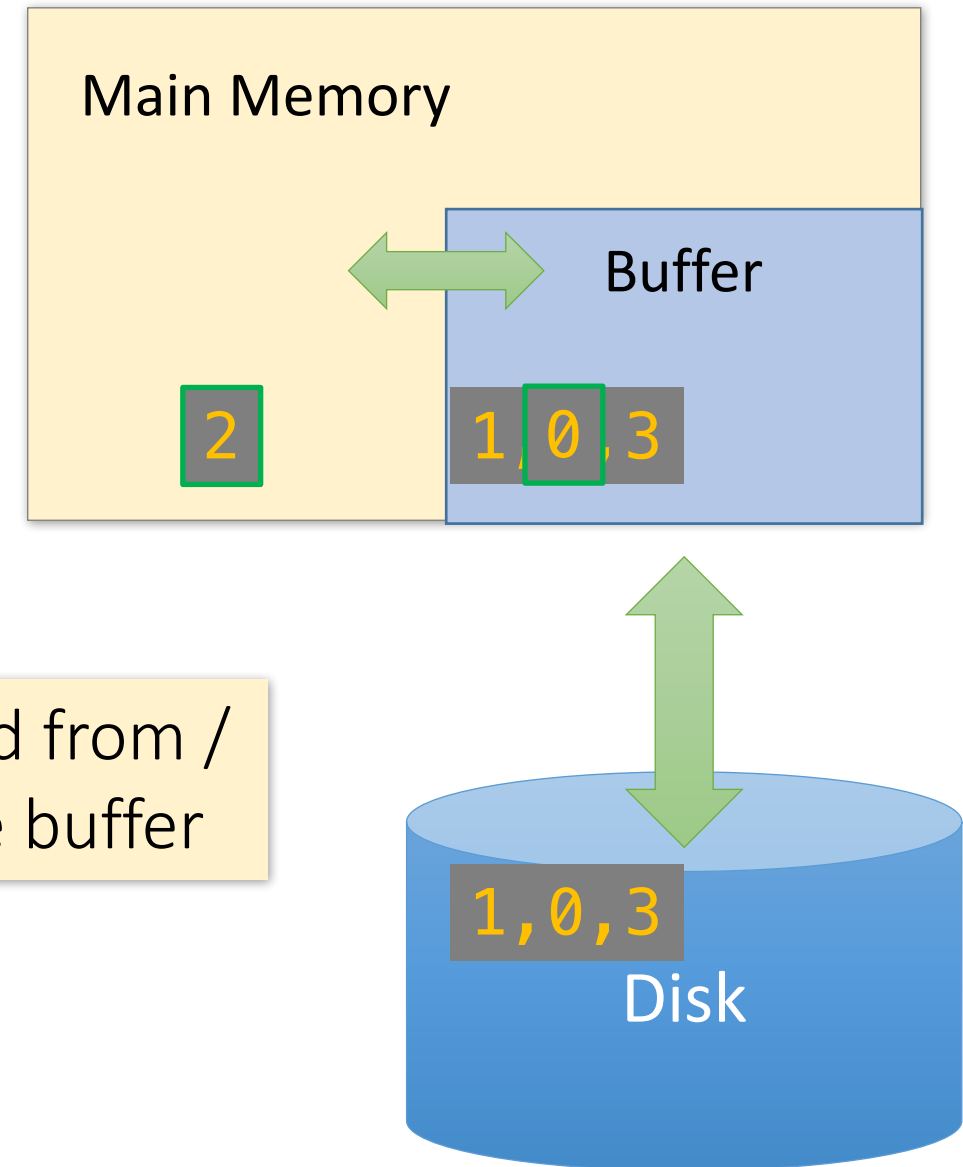
- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
- **Read(page)**: Read page from disk -> buffer *if not already in buffer*



The (Simplified) Buffer

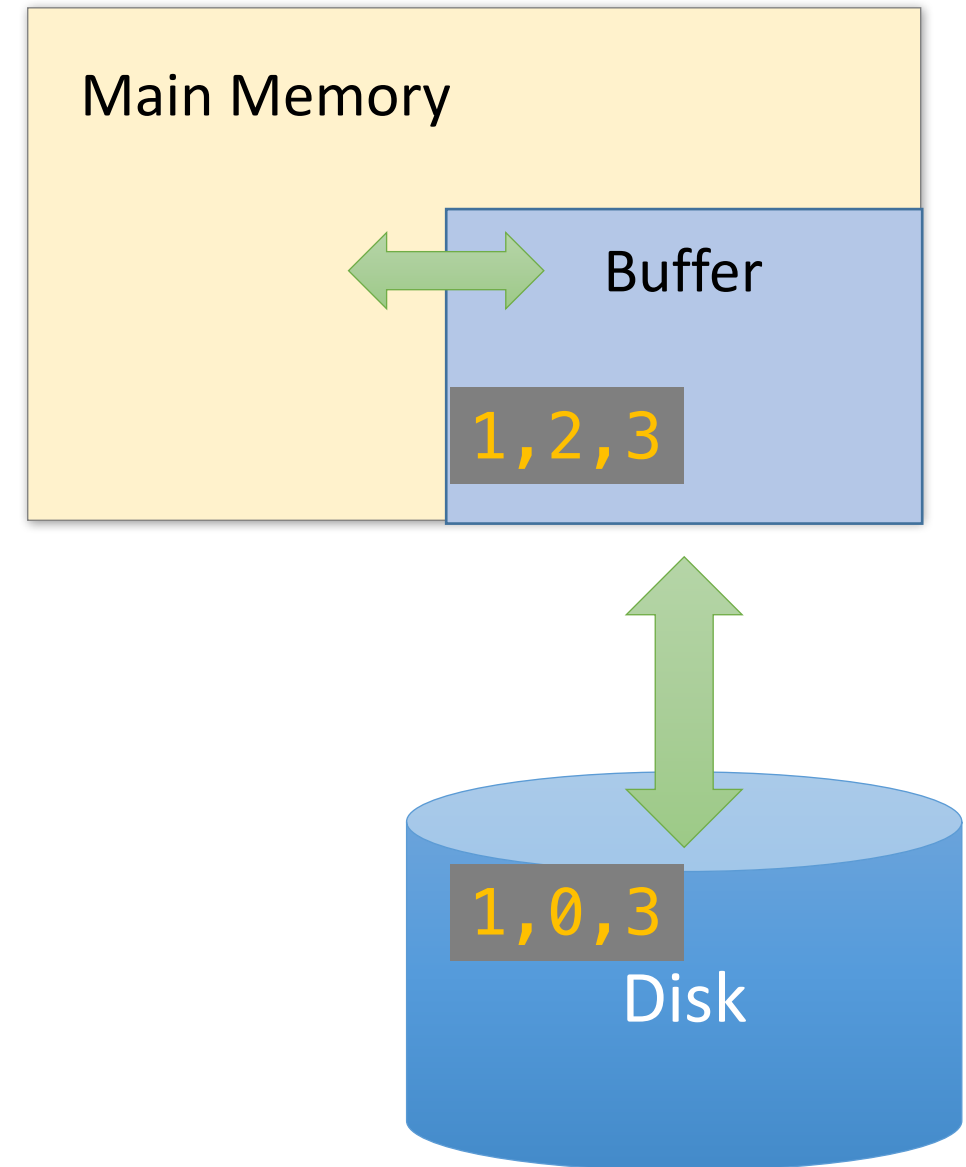
- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
- **Read(page)**: Read page from disk -> buffer *if not already in buffer*

Processes can then read from / write to the page in the buffer



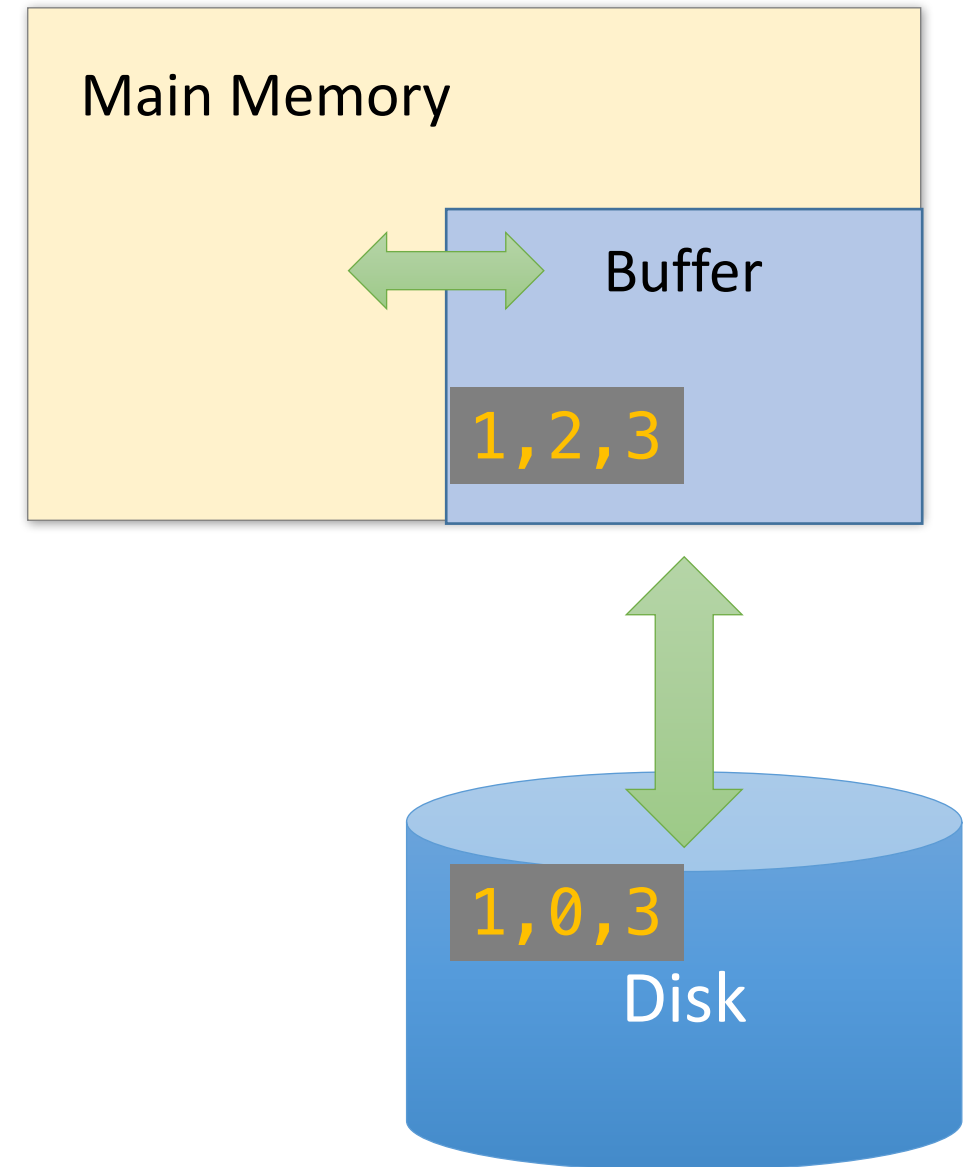
The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
 - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
 - **Flush(page)**: Evict page from buffer & write to disk



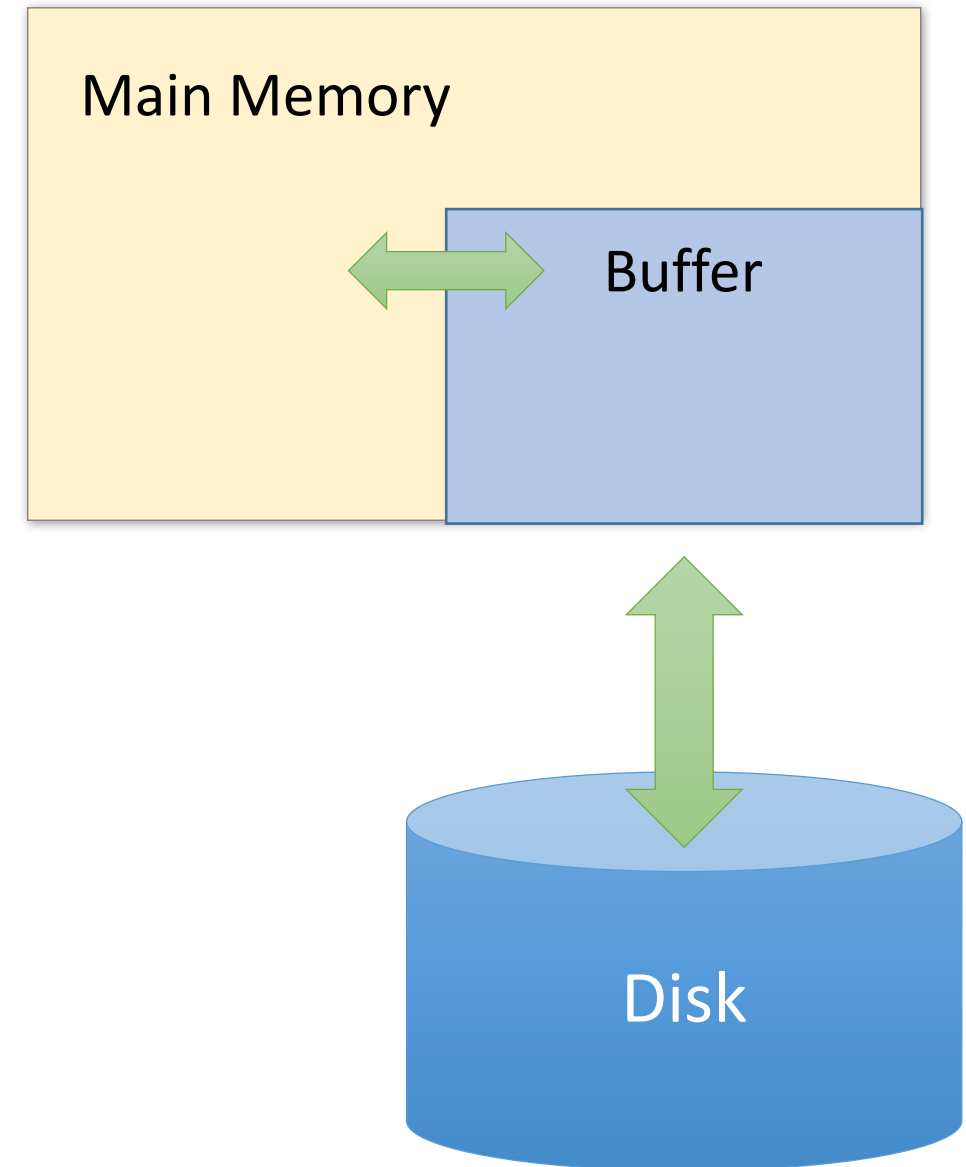
The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
 - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
 - **Flush(page)**: Evict page from buffer & write to disk
 - **Release(page)**: Evict page from buffer *without* writing to disk



Managing Disk: The DBMS Buffer

- Database maintains its own buffer
 - Why? The OS already does this...
 - DB knows more about access patterns.
 - Watch for how this shows up! (cf. *Sequential Flooding*)
 - Recovery and logging require ability to **flush** to disk.

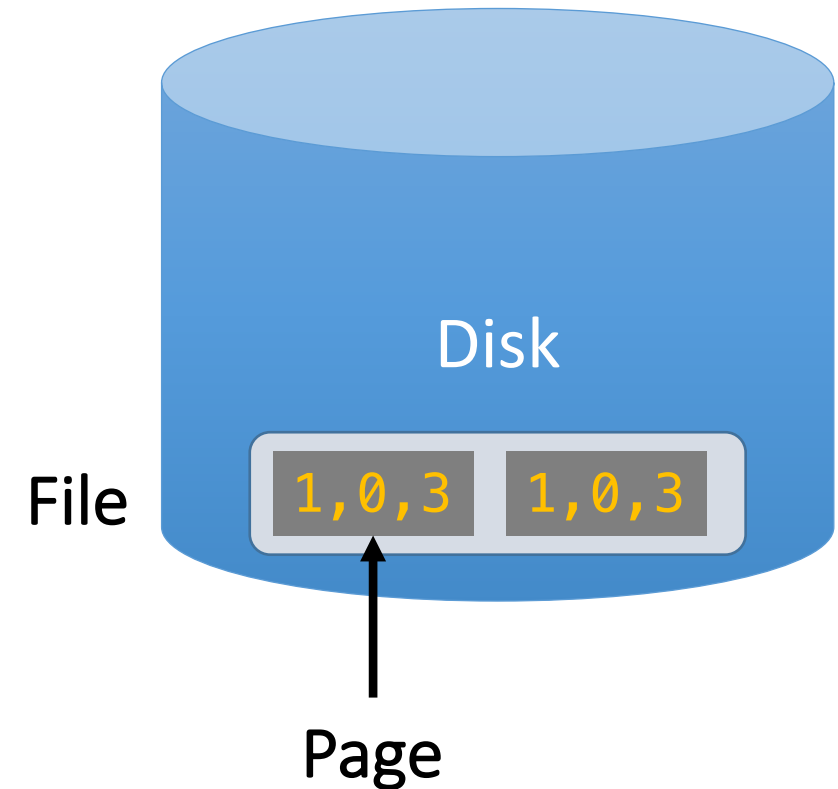


The Buffer Manager

- A **buffer manager** handles supporting operations for the buffer:
 - Primarily, handles & executes the “replacement policy”
 - i.e. finds a page in buffer to flush/release if buffer is full and a new page needs to be read in
- DBMSs typically implement their own buffer management routines

A Simplified Filesystem Model

- For us, a **page** is a ***fixed-sized array*** of memory
 - Think: One or more disk blocks
 - Interface:
 - write to an entry (called a **slot**) or set to “None”
 - DBMS also needs to handle variable length fields
 - Page layout is important for good hardware utilization as well (see 346)
- And a **file** is a ***variable-length list*** of pages
 - Interface: create / open / close; next_page(); etc.

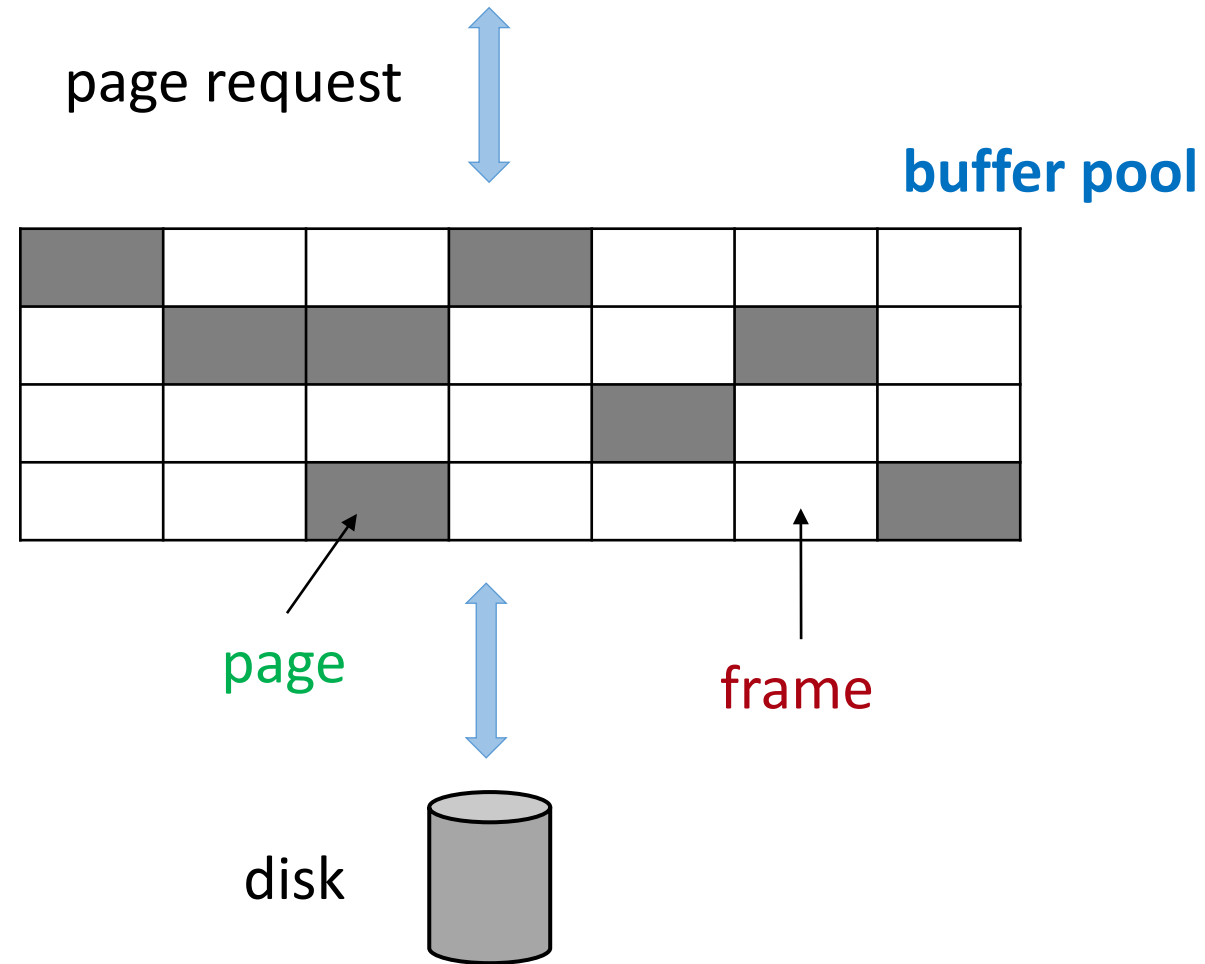


Buffer Manager

- Data must be in RAM for DBMS to operate on it
- All the pages may not fit into main memory

Buffer manager: responsible for bringing pages from disk to main memory as needed **pages** brought into main memory are in the **buffer pool** the buffer pool is partitioned into **frames**: slots for holding disk pages

Buffer Manager



Remember: Buffer Manager

- **Read(page)**: Read page from disk -> buffer *if not already in buffer*
- **Flush(page)**: Evict page from buffer & write to disk
- **Release(page)**: Evict page from buffer *without* writing to disk

Buffer replacement policy

- How do we choose a frame for replacement?
 - LRU (**L**east **R**ecently **U**sed)
 - Clock
 - MRU (**M**ost **R**ecently **U**sed)
 - FIFO, random, ...
- The replacement policy has big impact on # of I/O's (depends on the access pattern)
- To be continued!